

```
In [1]: # ECE 561 - Jackson Hellmers
import random
import numpy as np
import tensorflow as tf
import matplotlib.pyplot as plt
from matplotlib import pyplot as plt

# Load data and split into training and test sets
(x_train, y_train), (x_test, y_test) = tf.keras.datasets.mnist.load_data()

x_train = x_train.astype(float)
y_train = y_train.astype(float)
x_test = x_test.astype(float)
y_test = y_test.astype(float)

# note x_train is a numpy array with size (60000, 28 28)
print(np.shape(x_train))
print(np.shape(x_test))

# for this assignment, only use the first 2000 training examples
x_train = x_train[0:2000,:,:)
y_train = y_train[0:2000]


(60000, 28, 28)
(10000, 28, 28)
```

```
In [2]: # Prob. 1
# note that you can access the ith image in x_train as x_train[i-1,:,:)
# use plt.imshow(the_image, cmap=plt.cm.gray_r) to display image


img_18 = x_train[17,:,:)
plt.imshow(img_18, cmap=plt.cm.gray_r)

print("The image displayed is an '8'")
print("The label for the image displayed is",int(y_train[17]))
```


The image displayed is an '8'  
The label for the image displayed is 8




```
In [3]: # Prob. 2
for i in range(5):
    toShow = random.randint(0,len(y_train))
    plt.imshow(x_train[toShow,:,:), cmap=plt.cm.gray_r)
    plt.show()
    print("Label:",int(y_train[toShow]))
```




Label: 1




Label: 2



Label: 9



Label: 0



Label: 3

### Question 3

```
In [4]: ## Part A

## nearest neighbor classifier
def f(_x):
    assert np.shape(_x)==np.shape(x_train[0]) #check to make sure input is a single 28x28 np array
    distances = [np.linalg.norm(_x-x_train[i]) for i in range(len(y_train))] #calculate distance as pixel
    nearest_neighbor = np.argmin(distances) #find index of smallest distance
    return (y_train[nearest_neighbor], nearest_neighbor) #return the label of the nearest neighbor, and
```

```
In [5]: ## Part B

[label,index] = f(x_test[0])
correct = int(label)==int(y_test[0]) #check if prediction correct
misclass_loss = 0 if correct else 1 #calculate misclassification loss
norm_loss = np.linalg.norm(y_train[index]-y_test[0])**2 #calculate square error loss

print("Label of Actual and Predicted Match:",correct)
print("Predicted Label:",label)
print("Actual Label",y_test[0])
print("Misclassification Loss:",misclass_loss)
print("Squared Error:",norm_loss)
```

Label of Actual and Predicted Match: True  
 Predicted Label: 7.0  
 Actual Label 7.0  
 Misclassification Loss: 0  
 Squared Error: 0.0

```
In [6]: ## Part C

misclass_sum = 0
norm_sum = 0
incorrect = np.array([[[]],[[]]])
num_classifications = 1000
for i in range(num_classifications): #predict first 1000 elements of test data
    [label,index] = f(x_test[i])
    correct = int(label)==int(y_test[i])
    if not correct:
        misclass_sum += 1
        incorrect = np.append(incorrect,[[i],[index]],axis=1) #append the index of the incorrect prediction
    norm_sum += np.linalg.norm(y_train[index]-y_test[i])**2
misclass_emp = misclass_sum/num_classifications
norm_emp = norm_sum/num_classifications
print("Empirical Misclassification Loss",misclass_emp)
print("Empirical Square Error Loss",norm_emp)
```

Empirical Misclassification Loss 0.127  
 Empirical Square Error Loss 2.103

```
In [7]: ## Part D

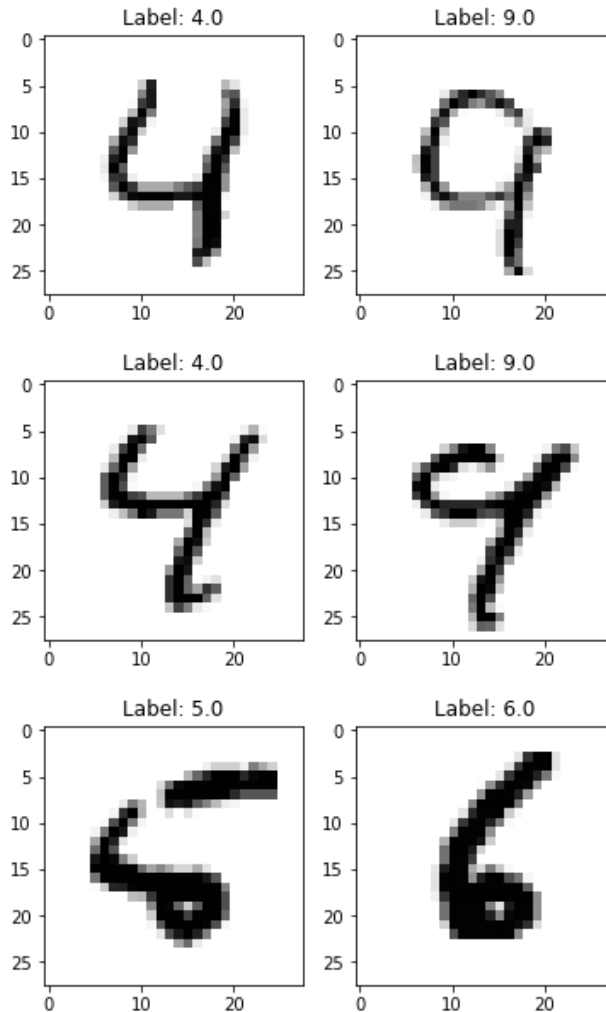
print("Actual Image on Left Nearest Neighbor on Right")
```

```

for j in range(3):
    actual = x_test[int(incorrect[0,j])]
    neighbor = x_train[int(incorrect[1,j])]
    f = plt.figure()
    f.add_subplot(1,2, 1)
    plt.imshow(actual, cmap=plt.cm.gray_r)
    plt.title("Label: "+str(y_test[int(incorrect[0,j])]))
    f.add_subplot(1,2, 2)
    plt.imshow(neighbor, cmap=plt.cm.gray_r)
    plt.title("Label: "+str(y_train[int(incorrect[1,j])]))
    plt.show(block=True)

```

Actual Image on Left Nearest Neighbor on Right



```

In [8]: # It is clear once visualizing the images how they were incorrectly predicted.
        # Many contain similar features and have a lot of overlap in size and curvature.
        # For some of the images I myself would have a hard time correctly labeling the image.

        # This exemplifies one of the draw backs of using nearest neighbors classification. As images
        # may have a large overlap in which pixels are on/off but contain different information.

```

## Question 4

```

In [9]: ## k nearest neighbors classifier
        def knn(_x,k):
            assert np.shape(_x)==np.shape(x_train[0]) #check to make sure input is a single 28x28 np array
            distances = [np.linalg.norm(_x-x_train[i]) for i in range(len(y_train))]
            distances = np.array(distances)
            distance_sorted = distances.argsort()
            distances_min = distance_sorted[0:k]
            return np.argmax(np.bincount(y_train[distances_min].astype(int))) #return the label of most occurring

```

```
In [10]: misclass_sum = 0
         square_error_sum = 0
         num_classifications = 1000
         k = 10
         for i in range(num_classifications):
             pred_label = knn(x_test[i],k)
             misclass_sum += 0 if (pred_label==int(y_test[i])) else 1
             square_error_sum += np.linalg.norm(pred_label-y_test[i])**2
         misclass_emp = misclass_sum/num_classifications
         square_error_emp = square_error_sum/num_classifications
         print("Empirical Misclassification Loss",misclass_emp)
         print("Square Error Loss",square_error_emp)
```

Empirical Misclassification Loss 0.136  
Square Error Loss 2.394

In [ ]: